

▼ Introduction to Big Data

Lecture 11. Classification

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import numpy as np
import pandas as pd
import seaborn as sns
```

▼ Classification with Personal Loan Data

- Experience
- Income
- Family
- CCAvg: Average monthly card spent
- Education: Education level (1: undergrad; 2, Graduate; 3; Advance)
- Mortgage
- Personal Loan: Personal Loan (1:Yes, 0:No)
- Securities account: Securities (1:Yes, 0:No)
- CD account: CD account (1:Yes, 0:No)
- Online: Online account (1:Yes, 0:No)
- CreditCard: Credit Card (1:Yes, 0:No)

```
PerLoan = pd.read_csv("/content/drive/MyDrive/[Lecture]/IntBigData/BigData_Python/11_Classification/personalLoan.csv")
PerLoan.head()
```

	ID	Age	Experience	Income	ZIP Code	Family	CCAvg	Education	Mortgage	Personal Loan	Securities Account	CD Account	Online	CreditCard
0	1	25	1	49	91107	4	1.6	1	0	0	1	0	0	0
1	2	45	19	34	90089	3	1.5	1	0	0	1	0	0	0
2	3	39	15	11	94720	1	1.0	1	0	0	0	0	0	0
3	4	35	9	100	94112	1	2.7	2	0	0	0	0	0	0



```
PerLoan.shape
```

```
(2500, 14)
```

```
PerLoan.columns
```

```
Index(['ID', 'Age', 'Experience', 'Income', 'ZIP Code', 'Family', 'CCAvg',
      'Education', 'Mortgage', 'Personal Loan', 'Securities Account',
      'CD Account', 'Online', 'CreditCard'],
      dtype='object')
```

```
PerLoan.rename(columns={'Personal Loan': 'PersonalLoan'}, inplace=True)
```

```
PerLoan.columns
```

```
Index(['ID', 'Age', 'Experience', 'Income', 'ZIP Code', 'Family', 'CCAvg',
      'Education', 'Mortgage', 'PersonalLoan', 'Securities Account',
      'CD Account', 'Online', 'CreditCard'],
      dtype='object')
```

```
PerLoan.describe()
```

	ID	Age	Experience	Income	ZIP Code	Family	CCAvg	Education	Mortgage	PersonalLoan	Securities Account
count	2500.00000	2500.000000	2500.000000	2500.0000	2500.000000	2500.00000	2500.000000	2500.00000	2500.000000	2500.000000	2500.00000
mean	1250.50000	45.346000	20.120800	74.4472	93135.691600	2.40800	1.951284	1.86560	57.388400	0.102400	0.109600
std	721.83216	11.519521	11.523824	46.6724	2420.763339	1.15986	1.795449	0.83655	100.816403	0.303234	0.312450

```
PerLoan.isnull().any()
```

```
ID                False
Age               False
Experience        False
Income           False
ZIP Code         False
Family           False
CCAvg            False
Education         False
Mortgage         False
PersonalLoan     False
Securities Account False
CD Account       False
Online           False
CreditCard       False
dtype: bool
```

```
PerLoan.count()
```

```
ID                2500
Age               2500
Experience        2500
Income           2500
ZIP Code         2500
Family           2500
CCAvg            2500
Education         2500
Mortgage         2500
PersonalLoan     2500
Securities Account 2500
CD Account       2500
Online           2500
CreditCard       2500
dtype: int64
```

```
PL_X = PerLoan[['Age', 'CCAvg', 'Income', 'Education']]
PL_Y = PerLoan['PersonalLoan']
```

▼ Logit Regression with statsmodels

```
from statsmodels.formula.api import logit

statsLogitModel = logit('PersonalLoan ~ Age + CCAvg + Income + Education', data=PerLoan)
statsLogitModel
```

<statsmodels.discrete.discrete_model.Logit at 0x7dbdf48699c0>

```
statsLogitModel_res = statsLogitModel.fit()
```

Optimization terminated successfully.
Current function value: 0.167135
Iterations 9

```
print(statsLogitModel_res.summary())
```

```

=====
                        Logit Regression Results
=====
Dep. Variable:          PersonalLoan   No. Observations:          2500
Model:                  Logit          Df Residuals:              2495
Method:                 MLE            Df Model:                  4
Date:                   Thu, 23 Nov 2023   Pseudo R-squ.:             0.4940
Time:                   13:14:44          Log-Likelihood:            -417.84
converged:              True             LL-Null:                   -825.81
Covariance Type:        nonrobust         LLR p-value:               2.695e-175
=====

```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-11.2330	0.694	-16.179	0.000	-12.594	-9.872
Age	0.0113	0.008	1.401	0.161	-0.005	0.027
CCAvg	0.1187	0.046	2.606	0.009	0.029	0.208
Income	0.0472	0.003	16.648	0.000	0.042	0.053
Education	1.6128	0.135	11.938	0.000	1.348	1.878

```
=====
```

```
statsLogitModel_res.params
```

```
Intercept    -11.232994
Age           0.011319
CCAvg        0.118712
Income       0.047228
Education    1.612840
dtype: float64
```

```
np.exp(statsLogitModel_res.params)
```

```
Intercept    0.000013
Age          1.011383
CCAvg        1.126045
Income       1.048361
Education    5.017038
dtype: float64
```

Logit Regression with sklearn

▼ Whole sample

```
from sklearn.linear_model import LogisticRegression
```

```
LogitModel0 = LogisticRegression()
```

```
LogitModel0_res = LogitModel0.fit(PL_X, PL_Y)
LogitModel0_res
```

```
▼ LogisticRegression
LogisticRegression()
```

```
LogitModel0_res.coef_
```

```
array([[0.01125087, 0.11722838, 0.04689883, 1.58410844]])
```

```
LogitModel0_res.intercept_
```

```
array([-11.12673563])
```

▼ Test and train sample



```
from sklearn.model_selection import train_test_split
```

```
PL_X_train, PL_X_test, PL_Y_train, PL_Y_test = train_test_split(PL_X, PL_Y, test_size=0.3,
                                                                random_state=0)
```

```
# Practice of Random Sampling
```



```
PL_X_train1, PL_X_test1, PL_Y_train1, PL_Y_test1 = train_test_split(PL_X, PL_Y,
                                                                    test_size=0.3)
```

```
PL_X_train1.head()
```

	Age	CCAvg	Income	Education	
2	39	1.0	11	1	
851	41	1.4	23	2	
2355	56	1.6	74	3	
360	35	1.3	55	1	
323	59	4.4	99	1	



```
PL_X_train1, PL_X_test1, PL_Y_train1, PL_Y_test1 = train_test_split(PL_X, PL_Y,
                                                                    test_size=0.3,
                                                                    random_state=1)
```

```
PL_X_train1.head()
```

	Age	CCAvg	Income	Education	
1181	25	0.2	65	1	
372	56	0.7	44	2	

```
PL_X_train1, PL_X_test1, PL_Y_train1, PL_Y_test1 = train_test_split(PL_X, PL_Y,
                                                                    test_size=0.3,
                                                                    random_state=1)
```


```
PL_X_train1.head()
```

	Age	CCAvg	Income	Education	
1181	25	0.2	65	1	
372	56	0.7	44	2	
137	49	0.4	128	1	
1831	47	0.4	30	2	
2296	27	0.2	82	1	

```
PL_X_train
```

	Age	CCAvg	Income	Education	
1988	52	0.3	18	1	

PL_X_test

	Age	CCAvg	Income	Education	
53	50	2.10	190	3	
2391	39	4.67	138	2	
2310	32	0.30	32	1	
728	45	4.40	114	2	
850	46	0.20	39	1	
...	
2285	48	2.40	114	3	
1165	43	1.70	113	1	
2438	62	0.30	29	3	
768	43	1.70	72	1	
381	55	2.30	73	3	

750 rows × 4 columns

PL_Y_train

```

1988    0
1969    0
1368    0
840     0
2214    0
...
1033    0
1731    1
763     0
835     0
1653    0

```

Name: PersonalLoan, Length: 1750, dtype: int64


```
PL_Y_test
```

```
53      1
2391    1
2310    0
728     1
850     0
```

```
..
2285    1
1165    0
2438    0
768     0
381     0
```

```
Name: PersonalLoan, Length: 750, dtype: int64
```

```
from sklearn.linear_model import LogisticRegression
```

```
LogitModel = LogisticRegression()
```

```
LogitModel.fit(PL_X_train, PL_Y_train)
```

```
▼ LogisticRegression
```

```
LogisticRegression()
```

```
LogitModel.coef_
```

```
array([[0.01166588, 0.07185039, 0.04623718, 1.44376302]])
```

```
LogitModel.intercept_
```

```
array([-10.5227233])
```

▼ Validation

```
PL_Y_pred = LogitModel.predict(PL_X)
```

```
PL_Y_train_pred = LogitModel.predict(PL_X_train)
```

```
PL_Y_test_pred = LogitModel.predict(PL_X_test)
```

```
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score, confusion_matrix
```

▼ Accuracy Score

```
accuracy_score(PL_Y_test, PL_Y_test_pred)
```

```
0.9546666666666667
```

▼ Recall Score

```
recall_score(PL_Y_test, PL_Y_test_pred)
```

```
0.65625
```

▼ Precision Score

```
precision_score(PL_Y_test, PL_Y_test_pred)
```

```
0.7777777777777778
```

▼ Specificity

```
tn, fp, fn, tp = confusion_matrix(PL_Y_test, PL_Y_test_pred).ravel()  
specificity = tn / (tn+fp)  
specificity
```

```
0.9825072886297376
```

